

Why AI Governance Must Be Enforced, Not Assumed

Enterprise Brief for CISO, Compliance, Risk, and AI Governance Leaders

LinkedIn Featured Asset #2

Core Position

AI governance is not established by policy documentation alone. A policy can define expectations, but it does not prove that an AI action was reviewed, authorized, traceable, or blocked when required evidence was missing.

For enterprise use, governance must operate at runtime. The control environment should validate policy, approval state, provenance, and audit evidence before an AI-assisted action is allowed to proceed.

What Runtime Enforcement Requires

Runtime enforcement means governance controls are evaluated before execution, not reconstructed after the fact. The system must be able to determine whether required approvals exist, whether the correct policy version applies, whether evidence is complete, and whether the action should be allowed or denied.

A governed AI workflow should treat missing validation as a failed validation. If approval, provenance, policy state, or audit evidence cannot be verified, the safe outcome is to block the action and preserve evidence of the denial.

Minimum Control Expectations

Control Area	Enterprise Requirement	Expected Evidence
Human Oversight	High-risk actions require explicit review before execution.	Reviewer identity, decision, timestamp, and approval scope.
Audit Evidence	Every allow or deny decision must be reconstructable.	Actor, policy version, validation result, and execution outcome.
Provenance	Inputs, model/workflow context, and decision path must be traceable.	Sustained workflow lineage, and evidence hashes where appropriate.
Fail-Closed Controls	Unknown or incomplete state must block execution.	Denial reason, missing-control indicator, and preserved audit trail.

Human Oversight Must Produce Evidence

Human oversight should not be limited to a dashboard view or informal review. It should create a durable decision record that identifies who reviewed the action, what they decided, when the decision occurred, and which policy context governed the decision.

This is especially important for high-risk workflows where security, compliance, legal, risk, or audit teams may need to verify that human judgment was applied before execution.

Provenance And Auditability Are Control Requirements

Provenance allows the organization to understand where an AI-assisted action came from and what evidence supported it. Auditability allows the organization to reconstruct the decision later without relying on memory, manual notes, or incomplete logs.

For enterprise AI governance, provenance and audit evidence should be generated as part of the decision path, not assembled after a review, incident, or regulator request.

Executive Takeaway

The central governance question is not whether an AI system is capable of producing an output. The enterprise question is whether the resulting action is authorized, policy-compliant, attributable, traceable, and defensible.

Governance that cannot enforce decisions at runtime remains an assumption. Governance that validates evidence before execution becomes a control.